

Bilinear Rank as Proof Cost: Fast Matrix-Multiplication Schemes and Field-Specific Flip Search for zkML

Greg Sidebottom

Research note — logic repo (<https://github.com/gsidebottom/logic>), matmul track, 2026-07-11. Companion to the 3×3 additive-complexity paper (`doc/matmul_adds_paper.md`, artifacts DOI 10.5281/zenodo.21240904), the rank-48 rigidity paper, and the flower note. The engine of §5 is `src/bin/flip23p.rs` in this repository.

Abstract

Zero-knowledge machine learning (zkML) — proving that a neural-network inference was computed correctly without re-executing or revealing it — is bottlenecked by the prover, whose work in circuit-based proof systems is counted in **multiplication constraints**: additions and constant-coefficient linear combinations are free. This is the exact inversion of the hardware cost model that keeps fast matrix multiplication out of GPUs, and it makes **bilinear rank** — the number of multiplications in a Strassen-like scheme — the direct cost driver for matrix products between witness values (private weights; attention). Three consequences. (1) Known schemes already pay: recursive rank-48 4×4 blocking proves a 768×768 witness $\times$ witness product with $\approx 4\times$ fewer constraints than the naive circuit, penalty-free because the schemes’ additions cost nothing. (2) Coefficient obstructions vanish: schemes over $\mathbb{Q}$ with dyadic (or any rational) coefficients are exact over a prime field, and numerical stability is not a concept. (3) Most interestingly, **rank is field-dependent** — rank 47 for 4×4 exists modulo 2 with no characteristic-0 equivalent known — and zkML fixes a short list of concrete fields (Goldilocks, BabyBear, M31, the BN254/BLS12-381 scalar fields) over which nobody has searched. We port our rational flip-graph engine to $\mathbb{F}_p$ (`flip23p`, Goldilocks first): over a prime field the engine is strictly stronger than its $\mathbb{Q}$ parent — no coefficient growth, no magnitude caps, and every coincidence-solving λ is admissible, so the coplanarity structure that rationality left mostly unusable becomes fully spendable. A verified rank-22 3×3 or rank-47 4×4 over a proof field would be a shippable constraint-count reduction with no analog in the characteristic-0 literature. We state the cost model precisely, including the honest carve-outs (sumcheck/GKR proves multiply matrices in $\sim n^2$ prover work and do not care about rank; public-weight linear layers cost no constraints), inventory which of our existing artifacts transfer, and report the port’s first measurements.

1 zkML in one page

A **zero-knowledge proof** (ZKP) lets a prover convince a verifier that a statement holds — here, “this output is what model M computes on input x ” — without the verifier re-executing the computation, and optionally without revealing x or M . Two families dominate. **SNARKs** (succinct non-interactive arguments of knowledge; Groth16 [6], PLONK [5]) give constant-or-polylog proof sizes over elliptic-curve-pairing fields such as the BN254 and BLS12-381 scalar fields (~ 254 -bit primes). **STARKs** [2] replace pairings with hashes, gain transparency and plausible post-quantum security, and run over small “FFT-friendly” primes chosen for fast number-theoretic transforms — Goldilocks $p = 2^{64} - 2^{32} + 1$ (Plonky2 [12]), BabyBear $p = 2^{31} - 2^{27} + 1$ (RISC Zero [13]), and M31 $= 2^{31} - 1$ (Circle STARKs [8]).

The asymmetry that defines the field: verification is cheap, **proving is 10^3 – $10^6\times$ the native computation**, all of it exact arithmetic in the proof field. zkML applications — proving a proprietary model was faithfully executed (model privacy), proving an inference inside a blockchain rollout, auditable AI where the checker is cheap — inherit that overhead on every multiply of every layer. Making matrix multiplication cheap *inside a proof* is therefore a different optimization problem from making it fast on silicon, with a different — in fact opposite — cost model.

2 What proving charges

The standard arithmetization, **R1CS** (rank-1 constraint systems, as in Groth16 [6]), expresses a computation as constraints of the form

$$\left(\sum_i a_i w_i\right) \times \left(\sum_i b_i w_i\right) = \left(\sum_i c_i w_i\right)$$

over the proof field, where w is the witness vector and the a_i, b_i, c_i are *constants*. Each constraint carries exactly one multiplication of witness values; the linear combinations inside the parentheses — additions, subtractions, scalings by arbitrary field constants — are free: they are folded into the constraint’s constant vectors and cost the prover nothing. PLONK-family systems [5] charge per gate/row with the same shape (custom gates and lookup arguments shift constants around but preserve the headline: **bilinear multiplications of witnesses are the metered resource**).

For an $n \times n$ product $C = A \cdot B$ where both operands are witnesses, each entry $C_{ij} = \sum_k A_{ik} B_{kj}$ is a sum of n witness×witness products; a sum of n products is not one rank-1 quadratic form — it takes n constraints. The naive circuit therefore costs n^3 constraints. Two situations make both operands witnesses in zkML: **private weights** (the commercially central case — the model owner proves inference without revealing the model) and **attention** ($QK^\top$ and scores $\cdot V$ multiply two activation matrices — witness×witness even when all weights are public). Conversely, a linear layer with *public* weights is a constant linear map and costs zero multiplication constraints; rank optimization has nothing to offer there.

A toy example, 2×2 over $\mathbb{F}_{17}$. Take witnesses $A = \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix}$, $B = \begin{pmatrix} 5 & 6 \\ 7 & 0 \end{pmatrix}$; then $C = A \cdot B \equiv \begin{pmatrix} 2 & 6 \\ 9 & 1 \end{pmatrix} \pmod{17}$. The naive circuit spends one constraint per product — eight constraints of the shape $u_1 = A_{11} \times B_{11}$ — and output sums like $C_{11} = u_1 + u_2$ cost nothing further (they are linear). Strassen’s rank-7 scheme spends **seven**: its first product is the single constraint

$$(w_{A_{11}} + w_{A_{22}}) \times (w_{B_{11}} + w_{B_{22}}) = w_{P_1},$$

where all four operand additions sit *inside* the constraint’s linear forms, free. There is exactly **one witness per proof**: a single vector w over $\mathbb{F}_{17}$ holding every wire value — a conventional leading 1 (which lets constraints use constants), the inputs, the seven products, the outputs. Concretely, all 20 coordinates:

$$w = (1 \mid \underbrace{1, 2, 3, 4}_A \mid \underbrace{5, 6, 7, 0}_B \mid \underbrace{8, 1, 6, 8, 0, 5, 3}_{P_1..P_7} \mid \underbrace{2, 6, 9, 1}_C).$$

The subscripted symbols above are coordinates of this one vector, and each constraint is checked by reading them off: for P_1 , $(1 + 4) \cdot (5 + 0) = 25 \equiv 8 \checkmark$, and likewise $P_2 = 7 \cdot 5 = 1$, $P_3 = 1 \cdot 6 = 6$, $P_4 = 4 \cdot 2 = 8$, $P_5 = 3 \cdot 0 = 0$, $P_6 = 2 \cdot 11 = 5$, $P_7 = (-2) \cdot 7 = 3$ (all mod 17). The free output combinations reproduce C exactly: $C_{11} = P_1 + P_4 - P_5 + P_7 = 19 \equiv 2$, $C_{12} = P_3 + P_5 = 6$, $C_{21} = P_2 + P_4 = 9$, $C_{22} = P_1 - P_2 + P_3 + P_6 = 18 \equiv 1$. (The prover commits to w and proves the constraints hold; the verifier never sees w — that is the “zero-knowledge”. Note the naive circuit’s witness carries eight product coordinates to Strassen’s seven, so lower rank also shortens the committed vector.) Seven constraints instead of eight — 12.5% at one level, compounding to $n^{2.807}$ under recursion — and dyadic coefficients would be equally at home: $\frac{1}{2}$ is just the constant 9 in $\mathbb{F}_{17}$ ($2 \cdot 9 = 18 \equiv 1$). This is §3’s story in miniature; the rank-23 and rank-48 schemes are the same picture with better ratios. (Appendix A runs a complete Setup $\rightarrow$ Prove $\rightarrow$ Verify on a two-constraint slice of this example — the whole SNARK pipeline at hand-checkable scale.)

3 Bilinear rank is constraint count

A bilinear scheme of rank r computes an $m \times m$ block product with r products of linear forms. In R1CS each product is exactly one constraint — the linear forms slide into the constraint’s free linear parts — so an $m \times m$ witness×witness block product costs r constraints instead of m^3 , and recursion compounds it: $T(n) = r T(n/m)$, i.e. $n^{\log_m r}$ constraints total, **with zero cost for the combination additions at every level**. The additions that price fast schemes out of silicon are literally free here; there is no numerical-stability tax because arithmetic is exact; and coefficients may be any field constants — dyadic $\frac{1}{2}$ ’s are simply the constant $(p + 1)/2 \pmod{p}$.

Idealized constraint counts for a witness $\times$ witness product (recursion to scalar base; real circuits stop at a base tile, which shifts absolute numbers but not the ordering):

scheme (exponent)	$n = 768$	vs naive	$n = 4096$	vs naive
naive (3.000)	4.53×10^8	1.0 $\times$	6.87×10^{10}	1.0 $\times$
$3 \times 3 : 23$ (2.854)	1.72×10^8	2.6 $\times$	2.05×10^{10}	3.4 $\times$
Strassen $2 \times 2 : 7$ (2.807)	1.26×10^8	3.6 $\times$	1.39×10^{10}	4.9 $\times$
$4 \times 4 : 48$ (2.792)	1.14×10^8	4.0$\times$	1.23×10^{10}	5.6$\times$
$3 \times 3 : 22$ (2.814), <i>if found</i>	1.32×10^8	3.4 $\times$	—	—
$4 \times 4 : 47$ (2.785), <i>if found</i>	1.09×10^8	4.2 $\times$	—	—

Two readings. First, **the wins in the top half are available today**: the rank-48 scheme [3, 10] with dyadic coefficients drops into any R1CS/PLONK matmul gadget as-is (our repository carries the verified scheme and its minimized linear networks). Second, the *marginal* rows show why new records matter here more than in silicon: each rank step is a few percent per recursion level, compounding — and unlike hardware, nothing eats the margin.

4 The honest carve-outs

Rank is **not** the lever everywhere:

- **Sumcheck/GKR provers do matmul in $\sim n^2$ prover field-ops** (Thaler [4]; GKR [7]); zkCNN [9] and successors build dedicated zkML provers on this line, and where such a prover applies, bilinear rank is irrelevant. The rank lever applies to the (large) circuit-based world: general-purpose SNARK toolchains, zkVMs executing matmul as arithmetic, PLONK/R1CS pipelines such as EZKL [11], and folding-based IVC (Nova [14]) whose per-step cost is the step circuit’s constraint count.
- **Public-weight linear layers are free**; rank helps witness $\times$ witness products only — private weights and attention.
- **Constraint count is not wall-clock**: prover time also includes commitments (multi-scalar multiplications / hashing); fewer constraints shrink those too, but constants differ by system.
- **Nonlinearities dominate some workloads** (range checks, lookups for ReLU/quantization); the matmul share is large in attention/convolution-heavy models, smaller in lookup-heavy quantized pipelines.

5 Field-specific rank: the open frontier, and flip23p

Bilinear rank depends on the coefficient field. The emblem: AlphaTensor’s rank-**47** 4×4 scheme exists over $\mathbb{F}_2$ [3], no characteristic-0 rank-47 is known, and our companion work proved the only $\mathbb{Q}$ -usable rank-48 class *rigid* under a large certified move system — evidence that characteristic 0 is genuinely constrained in ways a fixed finite field need not be. zkML nominates a short, concrete list of fields that matter commercially — Goldilocks, BabyBear, M31, BN254-Fr, BLS12-381-Fr — and, to our knowledge, **no rank search has ever been run over any of them**: the flip-graph literature works over $\mathbb{F}_2$ and small extensions [15], the explicit records over $\mathbb{Z}/\mathbb{Q}$. A verified 3×3 rank-22 — or 4×4 rank-47 — over Goldilocks would be a new kind of record with an immediate consumer: 2–4% fewer constraints per recursion level in deployed proof systems, compounding as in §3.

flip23p is our rational flip engine (splits, λ -flips, reductions, canonical gauge, exact Brent verification — the machinery of the companion papers) re-based onto $\mathbb{F}_p$, Goldilocks first. Over a prime field the engine is *strictly stronger* than its $\mathbb{Q}$ parent:

1. **No coefficient growth**. $\mathbb{Q}$ -walks need magnitude caps and dyadic bookkeeping; field elements don’t grow. Every cap in the $\mathbb{Q}$ engine is deleted, not loosened.

2. **Every solved flip is admissible.** Over $\mathbb{Q}$, a coincidence ($f_i + \lambda f_j \propto f_m$, the Cramer-solved move that makes flips productive) is usable only when λ is dyadic — the filter discarded most solutions. Over $\mathbb{F}_p$ every nondegenerate solution is a legal move: the coplanarity structure (“H4’s raw material” in the flower note) becomes fully spendable.
3. **The gauge simplifies:** monic normalization (leading coefficient 1, scalars folded into the c -slot); proportional $\Leftrightarrow$ equal, with none of $\mathbb{Q}$ ’s content/sign bookkeeping.

First measurements (2026-07-11; §8): the 55-addition record scheme loads and verifies over Goldilocks (729 Brent equations mod p); the census gate matches the $\mathbb{Q}$ engine exactly (9 shared pairs, weight 177, distinct-rows 40). One honest surprise: the seed’s solved-move silence is *not* the dyadic filter’s fault — even with all field λ available, its 9 shared pairs admit no coincidence solution; the coplanarity geometry genuinely fails there. Mobility arrives at once elsewhere: a 25-second smoke storm from the most mobile census seed executed 1.3M moves with 369K solved flips and $\sim 10K$ reductions — a far higher solved-flip fraction than the $\mathbb{Q}$ storm ever achieved, the freed λ paying immediately. No sub-23 rank has appeared in smoke tests (nor was one expected at that effort); the campaigns are the point.

6 What transfers from the existing program

- **Engines and discipline.** flip48/flip23’s storm, closure, census, novelty and thin-storm machinery ports mechanically (v1 of flip23p carries storm/census/closures; the beam-chase and exhaustive campaign modes follow). The operational rules hard-won in the companion papers — completion logs, full-frontier dumps, control baselines, budget caps — apply verbatim.
- **Seeds.** Every scheme we hold is a valid $\mathbb{F}_p$ scheme: the 17,376 database classes, our 53 lifted classes, the 46 storm-new classes, and the 467 mod-2-invisible dyadic $\mathbb{Q}$ -schemes (denominators are field constants). The mobility census that ranked $\mathbb{Q}$ -seeds re-runs unchanged.
- **Verified gadgets.** The Lean pipeline that certified the 55-add scheme extends naturally: a scheme-to-circuit emitter (circom / Halo2 templates for the 23- and 48-multiplication blocks) with a machine-checked correctness certificate is a deliverable the proof-systems world actually values.
- **What does not transfer:** the additive-complexity results (55 adds, the no-54 theorem). Additions are free here; that program’s value stays in classical hardware and in witness-generation speed.

7 Program

(1) Goldilocks campaigns: storm + closure portfolios over the census-ranked seeds; targets, in order of plausibility $\times$ value: any new-class harvest over $\mathbb{F}_p$, rank 22 at 3×3 , rank 47 at 4×4 (the 4×4 port is the same edit at 16 coordinates). (2) Small-field variants: BabyBear and M31 (31-bit arithmetic, $\sim 4\times$ faster walks), then BN254-Fr. (3) Rectangular tiles: transformer shapes want $\langle m, n, k \rangle$ records; the engine generalizes as flip48 $\rightarrow$ flip23 did. (4) Verified gadget emitter + benchmark note: constraint-count measurements in an EZKL-style stack, naive vs recursive-scheme, with Lean-certified gadget templates.

8 Reproduction

```
cargo build --release --bin flip23p
./target/release/flip23p --census --dir matmul/mm23      # gates
./target/release/flip23p --dir matmul/seeds23/13a5bd4n \
  --seconds 25 --threads 4 --out matmul/found23p        # smoke storm
# any verified rank <= 22 over Goldilocks lands in
# found23p/RECORDP_rank*.txt and is announced loudly
```

A Setup, Prove, Verify — a complete toy proof

This appendix runs the entire SNARK pipeline, Groth16-shaped, on a two-constraint slice of §2’s example, every number over $\mathbb{F}_{17}$. One honest simplification, flagged where it occurs: real systems wrap the values below in elliptic-curve encodings (“in the exponent”), which is what makes hiding cryptographic; the *arithmetic* — R1CS $\rightarrow$ QAP $\rightarrow$ divisibility check at a secret point — is shown exactly as real provers compute it.

The statement. For reference, the full toy matrices of §2, whose entries supply every private value below:

$$A = \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix}, \quad B = \begin{pmatrix} 5 & 6 \\ 7 & 0 \end{pmatrix}.$$

The prover claims: “I know private values $A_{11}, A_{22}, B_{11}, B_{12}, B_{22}$ (here 1, 4, 5, 6, 0 — the marked entries of A and B) such that $P_1 = (A_{11}+A_{22})(B_{11}+B_{22}) = 8$ and $P_3 = A_{11}(B_{12}-B_{22}) = 6$.” The claimed products (8, 6) are **public**; the five inputs stay private. The witness (public part first):

$$w = (1, P_1=8, P_3=6 \mid A_{11}=1, A_{22}=4, B_{11}=5, B_{12}=6, B_{22}=0)$$

and the two constraints (check: $5 \cdot 5 = 25 \equiv 8$, $1 \cdot 6 = 6$):

$$\begin{aligned} c_1 &: (w_{A_{11}} + w_{A_{22}}) \times (w_{B_{11}} + w_{B_{22}}) = w_{P_1} \\ c_2 &: (w_{A_{11}}) \times (w_{B_{12}} - w_{B_{22}}) = w_{P_3} \end{aligned}$$

Step 0 — constraints become one polynomial identity (the QAP). Assign constraint c_j to the point $x = j$, and interpolate each coordinate’s coefficients across constraints with the basis $L_1(x) = 2 - x$, $L_2(x) = x - 1$:

coord	A-side of	A-poly	B-side of	B-poly	C-side of	C-poly
A_{11}	c_1, c_2	1	—	0	—	0
A_{22}	c_1	$2 - x$	—	0	—	0
B_{11}	—	0	c_1	$2 - x$	—	0
B_{12}	—	0	c_2	$x - 1$	—	0
B_{22}	—	0	$c_1(+), c_2(-)$	$3 - 2x$	—	0
P_1	—	0	—	0	c_1	$2 - x$
P_3	—	0	—	0	c_2	$x - 1$

Weight each polynomial by its witness value and sum:

$$A(x) = 1 \cdot 1 + 4(2 - x) = 9 - 4x, \quad B(x) = 5(2 - x) + 6(x - 1) = x + 4, \quad C(x) = 8(2 - x) + 6(x - 1) = 10 - 2x.$$

Sanity: at $x=1$: $A \cdot B = 25 \equiv 8 = C$ (c_1 ✓); at $x=2$: $A \cdot B = 6 = C$ (c_2 ✓). **Both constraints hold** $\iff A(x)B(x) - C(x)$ **vanishes at** $x = 1, 2 \iff$ **it is divisible by** $Z(x) = (x - 1)(x - 2)$. Over $\mathbb{F}_{17}$:

$$P(x) = AB - C = -4x^2 - 5x + 26 \equiv 13x^2 + 12x + 9, \quad Z(x) = x^2 - 3x + 2, \quad H(x) = P/Z = 13$$

(exact division, remainder 0). The secret knowledge is compressed into: “I can exhibit witness-weighted A, B, C and a quotient H with $AB - C = HZ$.”

Step 1 — Setup (the trusted ceremony). Sample a secret point, say $\tau = 5$ — the “toxic waste” — and publish the powers of τ and the QAP polynomials evaluated at τ , each wrapped in an encoding that supports additions and scalings but hides the value (elliptic-curve points g^{τ^i} in real systems; bare numbers here, with hiding flagged). Then **destroy** τ : anyone holding it can forge proofs — hence the multi-party ceremonies around Groth16, and the transparent (no- τ) alternatives like STARKs. Published for the toy: $Z(5) = 4 \cdot 3 = 12$ and the encoded coordinate-polynomial evaluations at 5.

Step 2 — Prove. The prover, holding w , computes the *encoded* evaluations — each a **linear combination** of published encodings with witness coefficients (this is why R1CS’s free-linear-combination structure is the whole design):

$$A(5) = 9 - 20 \equiv 6, \quad B(5) = 9, \quad C(5) = 0, \quad H(5) = 13.$$

The proof is $\pi = (\text{enc}(A(5)), \text{enc}(B(5)), \text{enc}(C_{\text{priv}}(5)), \text{enc}(H(5)))$ — in Groth16, after optimizations, three curve points, ~ 200 bytes, *independent of circuit size*. The prover never learns τ .

Step 3 — Verify. The verifier reconstructs the *public* part of $C(\tau)$ itself from the claimed outputs $(8, 6)$ — the prover cannot lie about what is being claimed — then checks **one multiplicative relation** (one pairing equation in real systems; pairings are the tool that multiplies two hidden values exactly once):

$$A(\tau)B(\tau) - C(\tau) \stackrel{?}{=} H(\tau)Z(\tau) : \quad 6 \cdot 9 - 0 = 54 \equiv 3, \quad 13 \cdot 12 = 156 \equiv 3. \quad \checkmark$$

Why this convinces (soundness). A cheating prover without a valid witness needs $AB - C$ divisible by Z ; the best it can do is fake the relation at the single hidden point τ . Two distinct low-degree polynomials agree at a random point with probability $\leq \deg / |\mathbb{F}|$ — about $2/17$ in the toy (which is why $\mathbb{F}_{17}$ is a classroom field), about 2^{-250} over a real 254-bit field. Not knowing τ , the prover cannot aim: Schwartz-Zippel is the heart of the construction.

Why it reveals nothing (zero-knowledge). The verifier sees only encodings, never witness coordinates. (Full Groth16 additionally randomizes each proof with masking multiples of $Z(x)$ so two proofs of the same witness look independent; omitted for clarity.)

The connection to rank, one last time. Each constraint became one interpolation point, one row of the QAP table, one share of prover work and setup size. Prove a witness $\times$ witness matrix product with a rank- r scheme and this whole pipeline is r constraints per block instead of m^3 — the scheme’s additions never appear as constraints at any stage; they live inside the linear combinations, which the encodings support for free.

Acknowledgments

This work was carried out in an extended interactive collaboration with Claude (Anthropic; the Fable 5 and Opus 4.8 models), which implemented the engines and drafted this text under the author’s direction and review.

References

- [1] V. Strassen. *Gaussian elimination is not optimal*. Numer. Math. 13:354–356, 1969.
- [2] E. Ben-Sasson, I. Bentov, Y. Horesh, M. Riabzev. *Scalable, transparent, and post-quantum secure computational integrity*. ePrint 2018/046.
- [3] A. Fawzi et al. *Discovering faster matrix multiplication algorithms with reinforcement learning*. Nature 610:47–53, 2022.
- [4] J. Thaler. *Time-optimal interactive proofs for circuit evaluation*. CRYPTO 2013.
- [5] A. Gabizon, Z. Williamson, O. Ciobotaru. *PLONK*. ePrint 2019/953.
- [6] J. Groth. *On the size of pairing-based non-interactive arguments*. EUROCRYPT 2016.
- [7] S. Goldwasser, Y. T. Kalai, G. N. Rothblum. *Delegating computation: interactive proofs for muggles*. STOC 2008.
- [8] U. Haböck, D. Levit, S. Papini. *Circle STARKs*. ePrint 2024/278.

- [9] T. Liu, X. Xie, Y. Zhang. *zkCNN*. CCS 2021.
- [10] J.-G. Dumas, C. Pernet, A. Sedoglavic. arXiv:2506.13242; A. Novikov et al. *AlphaEvolve*. arXiv:2506.13131.
- [11] EZKL. github.com/zkonduit/ezkl.
- [12] Polygon Zero. *Plonky2*. Whitepaper, 2022.
- [13] RISC Zero. *zkVM documentation*, 2023.
- [14] A. Kothapalli, S. Setty, I. Tzialla. *Nova*. CRYPTO 2022.
- [15] M. Kauers, J. Moosbauer. *Flip graphs for matrix multiplication*. arXiv:2212.01175.
- [16] G. Grassi et al. *Poseidon*. USENIX Security 2021.
- [17] G. Sidebottom. *A 55-addition rank-23 scheme for 3×3 matrix multiplication*. This repository, 2026. DOI 10.5281/zenodo.21240904.
- [18] G. Sidebottom. *Rigidity of the rank-48 4×4 scheme and The Flower*. Companion notes, this repository, 2026.