

The Rank-21 Program: Reproduction, Soundness, Exhaustive Negatives, and a Certified-Search Architecture for $\langle 3, 3, 3 \rangle$ over $\mathbb{F}_2$

Greg Sidebottom

August 31, 2026 (working note); postscripts September 2, 4 and 6, 2026

Abstract

We report on a program aimed at improving the lower bound for the rank of the 3×3 matrix multiplication tensor over $\mathbb{F}_2$, currently $20 \leq R_{\mathbb{F}_2}(\langle 3, 3, 3 \rangle) \leq 23$ (Wang’s automated lower bound; Laderman’s scheme reduced mod 2). The contributions are: (i) an independent reproduction and strengthening of Wang’s certified DP-over-restrictions bound, including budget-scaled lifts of his orbit lattice; (ii) a sound extension of his framework to *mixed-side* substitution, including an activity guard whose necessity we demonstrate by measurement, and a family of exhaustive negative results showing that substitution with flattening-class leaves cannot exceed the lattice values; (iii) an exact determination of our substitution game’s value (18) under successively stronger rule sets, with machine-checked certificates; (iv) a certified-search (CDCL(T)-style) architecture over scheme space whose propagator and symmetry components each deliver their theoretical strength (measured ablations of $746\times$ and $632,384\times$), reaching an independent, exhaustively-searched $R_{\mathbb{F}_2} > 15$ in 3.25 hours (root set completed in the postscript); and (v) an empirical structure theorem for the adversary’s dangerous moves in the substitution game — *the Rank Law* — identifying them with rank-one covectors, which yields a concrete, well-posed lemma — proved by orbit exhaustion in the postscript, together with its level-2 form. We close with a convergence thesis: three structurally different attacks terminate at the same missing object, a cheap rank lower-bound oracle of strength 15–19 for arbitrary tensors, and we locate the genuine frontier at scheme size 20, where every known bound is structurally one short of what novelty requires.

1 Setting and known bounds

Let $\mathcal{T} = \langle 3, 3, 3 \rangle \in \mathbb{F}_2^9 \otimes \mathbb{F}_2^9 \otimes \mathbb{F}_2^9$ denote the matrix multiplication tensor in the trace convention $\mathcal{T} = \sum_{i,j,k} x_{ij} \otimes y_{jk} \otimes z_{ki}$, a 729-bit object with 27 nonzero entries. Its rank $R_{\mathbb{F}_2}(\mathcal{T})$ is the minimum number of products $\alpha \otimes \beta \otimes \gamma$ (rank-one tensors) whose XOR is $\mathcal{T}$; schemes of r products are exactly solutions of the Brent equations with r multiplications. Known bounds: $R_{\mathbb{F}_2}(\mathcal{T}) \geq 20$ (Wang, arXiv:2603.07280, with a machine-verifiable certificate of ~ 32 MiB), and $R_{\mathbb{F}_2}(\mathcal{T}) \leq 23$ (Laderman’s scheme reduced mod 2; independently, our 55-addition rank-23 scheme). Bläser’s ≥ 19 holds over every field. The program’s goal is 21.

Throughout, *flattening* bounds are matrix ranks of the three unfoldings (≤ 9 here); K_p denotes the Koszul flattening with wedge parameter p (division bound $\lceil \text{rank} / \binom{8}{p} \rceil$); and the *substitution lemma* is used in the form: for an active direction (one the tensor depends on), $R(T) \geq 1 + \min_{\lambda} R(T|_{\lambda})$ where λ ranges over all foldings of that direction.

A structural datum that organizes everything below: rank methods — bounds expressible as matrix ranks of linearly-constructed flattenings, including Koszul, Young and Strassen’s equa-

tions — are barrier-limited near $2n^2 = 18$ for matmul-type tensors. Our measurements respect this barrier with uncanny precision: flattenings reach 9, Strassen’s commutator bound 12–13, Koszul 14–16, our exhaustive game with such leaves exactly 18. Wang’s 20 lives beyond the barrier because substitution-DP is *recursive amplification*, not a rank method. This is the frame in which “new leaf mathematics” must be read (§6, §10).

2 Reproducing and strengthening the lattice DP

We reproduced Wang’s result locally from his repository (verifier included), then strengthened his certificate by re-running his backtracking technique at $30\times$ the shipped step budget over the whole restriction lattice. The wave lifted every layer — e.g. the fourteen codimension-2 orbits from $\{17\times 1, 18\times 13\}$ to $\{18\times 6, 19\times 8\}$ — but the root held at 20: a diminishing-returns profile, quantified. (A root of 21 requires, among more, all fourteen codim-2 orbits at ≥ 19 ; six remain at 18 and resisted 3×10^9 steps each.) Two engineering contributions accompany the reproduction: a diagnosis and fix of an allocator-retention memory ratchet in long runs (inline DFS-cache keys; heap pressure relief), and per-orbit byte-level certificate merging. All lifted certificates re-verify under his verifier.

3 Mixed-side substitution: soundness and negatives

Wang’s framework restricts only the *A*-side. Motivated by our game measurements (*A*-only value 16 vs. mixed value 18; §4), we implemented the *fifth technique*: chains mixing *A*- and *B*-side substitutions inside his backtracking DFS, with a backward-compatible proof format (per-record leaf kinds, per-node side bits) and a mirrored verifier.

The activity guard. The naive mixed-leaf bound (chain length plus flattening of the doubly-restricted tensor) is *unsound*: pure-*A* chains get activity for free (*A*-flattening injectivity survives *A*-restriction), but under double-sided restriction a substitution can kill a direction the tensor no longer depends on, inflating the bound by +1 per inactive step. The sound credit is rank-based: $k_B = \dim B - \dim(W + D)$, where W is the kernel of the *B*-rows and D the dead *B*-subspace of the *A*-restricted tensor. We validated the formula on engineered deactivations (independent Python implementation and a C++ probe), then measured it in production: per-orbit docked credit ranged from 0 to 8.6×10^9 — the guard is not vacuous, and the unguarded rule’s 99.99% “rescue rates” were partly phantom credits.

Exhaustive negatives. With soundness in place, equal-budget experiments returned a clean null: guarded mixed play lifted no codim-2 orbit at 3×10^8 steps (nor did *A*-only). Escalating down-lattice where exhaustion is affordable produced eight decisive verdicts, *all FAILS*: at codim-4 and codim-3, mixed substitution with table leaves cannot lift Wang’s values by one, exhaustively. A three-sided variant (*C*-moves) inflates the game ($\geq 6.6\times$ nodes on refutations that were 16–40 minutes under two sides) with no observed proving power. The binding constraint is identified and measured: once play leaves the pure-*A* lattice, every available leaf collapses to flattening class (≈ 14 –16), and the adversary grinds targets down against them.

4 The exhaustive substitution game

Independently of Wang’s framework, we built an exact game engine over (U_A, U_B, U_C) restriction states with sandwich-group canonicalization, stabilizer-reduced adversary fanouts, and memoized decision ladders. Its verdicts, each with machine-checked certificates replayed by an independent Python checker (including isomorphism edges): game value 17, then 18 under support/spanning refinements and rank-profile case splits; three exhaustive FAILS at 19 under every rule set tried. A methodological lesson is recorded with the mathematics: an early version “proved” $R_{\mathbb{F}_2}(\langle 2, 2, 2 \rangle) \geq 12$ (true value 7) because kill moves scored on the zero tensor *and the checker re-encoded the same rule*. Certificate checkers validate transcription, not semantics; every new proof rule since has been gated on known exact values first.

5 Certified scheme search (CDCL(T) over products)

The constructive dual: prove $R_{\mathbb{F}_2} > r$ by exhausting r -product schemes, with theory propagators supplying the reasoning that resolution lacks. Decisions are whole products; the residual (tensor XOR chosen products) carries native parity; propagators prune by residual rank bounds; symmetry enters as constructive orbit enumeration.

Gates at 2×2 . The stage-1 prototype reproduces both certified values: $r = 6$ UNSAT (matching our hydra_satsuma DRAT-verified result) and $r = 7$ SAT with an independently verified Strassen-class witness. The measured ablation: one ply of sound substitution reasoning cut the $r = 6$ tree from 3.22×10^9 to 4.3×10^6 nodes (746 $\times$).

Scaling to 3×3 . Lazy enumeration over the $511^3 \approx 1.3 \times 10^8$ products; triple-layout residuals making every flattening nine contiguous 81-bit extracts; constructive first-product canonicalization under the sandwich group ($133,432,831 \rightarrow 211$ orbit representatives; the level-1 frontier *is* the orbit count); level-2 stabilizer minimality (median 16 $\times$); a Strassen-commutator propagator (strength 12–13 at 10–30 μ s, validated by a constructed-rank gate: tensors built as XOR of k products have rank $\leq k$ by construction); Koszul at the shallow levels. The ladder:

rung	verdict	nodes	wall
$r = 9$	UNSAT (flatten)	783,363 $\rightarrow$ 211	0.1 s
$r = 10$	UNSAT (probe)	211	0.1 s
$r = 11$ –14	UNSAT (Koszul at level 1)	211 each	0.1 s each
$r = 15$	UNSAT (killer-directed deep probe)	211	3.25 h
$r = 16$	out of reach	—	≥ 70 core-h on one root, unresolved

The $r = 15$ rung is the architecture’s current independent reach, $R_{\mathbb{F}_2} > 15$: two rungs past the Koszul ceiling, taken by *lazy adversary-directed recursion* — at each node a cheap-floor chain (flatten $\rightarrow$ Strassen $\rightarrow$ Koszul) settles most folds, and only the *killers* (Koszul-dropping folds) recurse, with targets falling into Strassen range within two levels. Fixed-sweep plies, by contrast, were measured intractable (16.7 hours resolving fewer than 8 of 211 roots). Engineering lessons are recorded honestly: Method-of-Four-Russians elimination, equality-gated, was *benched out* (dense wins do not transfer to the sparse wedge shape); a leading-bit rewrite of the 9×81 rank

kernel gave a universal $2.1\times$; five unsound artifacts across the program (an α -canonical cut, a $\lambda=0$ -only probe, side-slot pollution, activity-free mixed leaves, a row-reversed matrix inverse) were all caught by known-value or constructed-value gates before contaminating any claim.

6 The Rank Law

Instrumenting the $r = 15$ probes produced a dataset of every Koszul-dropping fold at the 211 roots (10,338 killers; the probe's pivot is always the cell $(2, 2)$; every drop is $14 \rightarrow 13$ exactly). Writing the folding covector as $\varphi = e_{(2,2)} + \lambda$ and viewing it as a 3×3 matrix over $\mathbb{F}_2$:

Observation 1 (Rank Law; 211 roots $\times$ 3 sides = 633 pairs).

1. $\text{rank}(\varphi) = 1 \implies$ *killer, with zero exceptions. These are the $4 \times 4 = 16$ rectangle indicators through the pivot cell — exactly the observed sixteen killers per pair.*
2. $\text{rank}(\varphi) = 3 \implies$ *non-killer, with zero exceptions.*
3. $\text{rank}(\varphi) = 2$: *non-killer at 595 of 633 pairs; killer only at 38 exceptional pairs concentrated on 30 roots (210 extra folds).*

The reading: the adversary's dangerous moves are *product-shaped* — rank-one covectors aligned with the tensor's own rank-one structure, the same objects our game's rank-profile splits singled out. The law immediately poses a well-defined lemma:

Conjecture 1 (case-split lemma). *On the level-1 residuals (matmul minus one product, restricted as in the probe), folding by a covector of matrix rank 3 preserves $K_4 \geq 14$.*

Proven — and extended across the rank-2 boundary, whose 38 exceptions appear tied to low-rank structure of the root representatives — the lemma removes 168–240 of every 256 probe evaluations, repricing the $r = 16$ rung from ≥ 1.5 core-years to days. Independent of that payoff, the law is the first structural characterization of the substitution game's critical directions that we are aware of, and the killer set's algebraic description (points of an affine determinantal pencil over $\mathbb{F}_2^8$) suggests it is provable. (*It is: see Section 7, where the lemma and its level-2 form are proved by orbit exhaustion, and where the repricing estimate above is corrected by measurement.*)

7 Postscript (September 2, 2026): the lemma proved, a symmetry defect repaired, and the $r = 16$ wall quantified

7.1 The case-split lemma, by orbit exhaustion

Fix a fold vector $v \in A = M_3(\mathbb{F}_2)$ (the covector of Section 6 is $\text{tr}(\cdot v)$) and a rank-one $m = \alpha \otimes \beta \otimes \gamma$. The Koszul data of $(\mathcal{T} + m)|_{\ker v}$ is a sandwich invariant of the pair (m, v) , and the sandwich group is transitive on the rank classes of v ; so it suffices to fix v at a class representative and to enumerate the orbits of m under $\text{Stab}(v) \times GL_3$. We enumerate them exactly by a staged lexicographically-minimal canonical form and require the orbit sizes to sum to $511^3 = 133,432,831$; every orbit is then evaluated.

Theorem 1 (case-split lemma, level 1). *For every rank-one m over $\mathbb{F}_2$ and every v of matrix rank 3, on any side, $K_{\leq 4}((\mathcal{T} + m)|_{\ker v}) = 14$. The side- A , $p = 4$ wedge rank never drops below*

912 (threshold 911); over the 7,029 orbits the drop is 0 (6,505 orbits, 122,193,407 products) or exactly 6 (524 orbits, 11,239,424 products), each 6 split 3+3 across the two honest 8-dimensional wedges $p = 3, 4$ on A/v (both of rank 459).

Theorem 2 (rank-2 boundary and rank-1 inheritance). *For v of rank 2 (14,070 orbits) the fold kills on exactly 108 orbits (774,144 products, 0.58%), each dropping the wedge by exactly $12 = 6 + 6$ to Koszul 13; the boundary is a finite table. For v of rank 1 (3,556 orbits) the base value is 900 and every product is a killer; the perturbation can raise the wedge rank by 6, and the 6-quantization breaks ($+22 = 10 + 12$) only where it no longer matters.*

Gates: the encoded action fixes $\mathcal{T}$ term-by-term (random elements); random (m, v) transports agree; every (root, side, λ) fold of the $r = 15$ probes is transported to its representative and compared with the direct computation (162,048 triples, zero mismatches), and the predicted killer set of the Rank Law’s roots reproduces the 10,338-line dump byte for byte. The necessary conditions for the 108 rank-2 killers are $\text{tr}(\alpha\beta\gamma) = 1$, $\text{tr}(v\beta\gamma) = 0$, $\text{rank } \alpha \leq 2$, $\alpha \neq v$; no combination of up to three matrix-rank/trace invariants is a pure classifier, so the structural (isotypic) proof stays open — the exhaustion is the proof.

Theorem 3 (level 2). *Two successive side- A folds quotient by a 2-dimensional $U \subset A$. The 43,435 such U fall into 17 pencil classes. Every class whose pencil contains a rank-one matrix, and the $(2, 2, 2)$ pencil with stabilizer of order 288, has $K(\mathcal{T}|_{\ker U}) \in \{12, 13\}$ and admits no leaf for any m (the perturbation raises the wedge by at most 12, short of 911). The remaining eight classes keep Koszul 14 (wedge 918, the level-1 value) on 96.65–99.90% of products, the killers dropping by exactly 12 or 18. Over all (U, m) pairs the level-2 leaf fraction is 71.92% (1.65 million orbits evaluated; 300 random transports agree).*

7.2 A symmetry defect and its repair

The invariance gate of the lemma tool failed on first run: the scheme search’s C-side action, $\gamma \mapsto P^{-1}\gamma R^{-1}$, had been ported verbatim from the game engine, which indexes C as (i, k) ; the search indexes C as (k, i) , where the same group element acts as $R^{-T}\gamma P^{-T}$. The ported action is *not* a symmetry of $\mathcal{T}$, so the “211 first-product orbit representatives” of Section 5 were representatives of a fake partition: an audit under the correct action shows they meet 159 of the 211 true orbits, leaving 52 never probed. (The α and β stages were correct; the counts coincide because the defect is a bijection of the orbit set.) Repair: $r = 14$ re-exhausts on the corrected representatives in 0.7 s; the killer-directed deep probe proves $\text{rank} \geq 15$ on all 52 missed orbits in 394 s. The $r = 15$ row of the ladder now rests on a complete, gate-verified representative set. The Rank Law’s statements concern the dumped residuals themselves and stand; the lemma above quantifies over all products. The lesson belongs with the method acknowledgment: an orbit reduction must be gated on the object it claims to preserve, not on a comment saying it was verified elsewhere.

7.3 The $r = 16$ wall, quantified

With the lemma in hand we rebuilt the probe for the rung: a work-queue implementation that holds all twelve cores (the fanned and root-parallel modes tail to eight and one), removal of *no-op folds* (a fold into a zero slice is a relabeling, not a quotient; a minimal decomposition projects off any zero slice, so half of level-2 and three quarters of level-3 evaluations were wasted), fold-orbit pruning under the root’s stabilizer (one representative per orbit proves the whole conjunction,

since the rank bound is an isomorphism invariant), and a transposition table keyed by the reduced bases of the folded subspaces (an exact key: the residual is the projection along U onto the non-pivot coordinates, unique given U). Verdicts were gated against the serial probe at every step. Cumulatively, the $r = 15$ proof of one root went from 65,792 to 1,473 tasks, and the hardest root from 42.8 to 13.3 seconds. The $r = 16$ rung on root 0 nevertheless stayed unresolved at every attempt (30 minutes on 12 cores each; the August 31 attempt had reached 70 core-hours). The per-depth profile explains why (nodes evaluated after 29 minutes, memo run):

depth (target)	4 (12)	5 (11)	6 (10)	7 (9)	8 (8)
nodes evaluated	2,674	87,050	199,245	846,677	2,825,251
Koszul leaves	14	78,414	143,622	413,355	2,564,972
genuine failures	0	0	0	0	181,398

Pure side- A chains bottom out at depth 7–8 on pencils all three of whose combinations have rank ≤ 7 ; the probe then switches sides, and the mixed keys (seven A -folds and one B -fold: $2^{14} \times 256$) are millions, so the deep layers are side-switch fan-outs, not revisits (memo hit rate 15–50%). Balanced side selection (fold the side with most active coordinates) is verdict-equivalent but loses at $r = 15$ by two orders of magnitude, because it keeps every level’s fan at 256. In 29 minutes the search had completed 36 of the roughly 3,000 depth-3 nodes: the tree is $\sim 100\times$ the budget with every lever applied.

Two oracle routes for the small residuals were tested. Plain SAT (cadical on a Tseitin encoding of “rank $\leq r$ ”) is dead: a 9×9 pencil of known rank 9 needs 3.5 s for the satisfiable side and exceeds 60 s to refute rank 8; random pencils time out at every $r \in [8, 11]$. An exact pencil oracle via the Kronecker form is feasible but the textbook count is wrong over $\mathbb{F}_2$: brute force ($\min_X \text{rank } X + \text{rank}(I+X) + \text{rank}(A+X)$, $n \leq 4$) shows an irreducible cubic block contributes +2, not +1, as do $(x+1)^4$ and irreducible quartics; block additivity held in every sample. Such an oracle would settle depth 7 exactly and remove the depth-8 layer — about 70% of the work, a $3\times$ — and would not reach the rung. The wall is depths 5–6: 3- and 4-slice residuals that must be proved to have rank 10–11, where Koszul caps at 13 but fails on the hard instances and no cheap exact method is in sight.

Standing verdict (superseded the same day). The paragraph above was written before the frontier was examined directly. Dumping the depth-5/6 residuals showed them to be *concise* $3 \times 6 \times 6$ and $4 \times 6 \times 6$ tensors (common kernels of dimension 3, row and column spans 6) of rank exactly the target — true nodes that Koszul misses by one — being folded in the original 9×9 coordinates (fans of 2^8 where 2^5 suffices) and evaluated with a 9-dimensional wedge (milliseconds where microseconds suffice), under a depth cap that cut off the small subtrees that finish such proofs by flatten. Rewriting every residual in bases of its slice spans (an isomorphism), keying the memo by the concise tensor, trimming the Koszul wedge to the concise dimensions and allowing depth 12 — each step gated on known values ($\langle 2, 2, 2 \rangle$ proves 7 and refuses 8; $\langle 2, 2, 3 \rangle$ refuses 12) and on the recorded $r = 15$ verdicts — reduced the hardest $r = 15$ root from 42.8 s to 0.18 s and resolved root 0 at $r = 16$ in 4.6 s.

Theorem 4 ($r = 16$ rung, independent). $R_{\mathbb{F}_2}(\langle 3, 3, 3 \rangle) > 16$: the killer-directed deep probe at target 16 proves $\text{rank}(\mathcal{T} + m) \geq 16$ on all 211 true first-product orbit representatives (3,893 s on 12 cores; median 9.6 s per root, maximum 124.6 s; 3.68×10^8 tasks).

The lesson is methodological and belongs with the acknowledgment below: when a search explodes, dump its frontier and look at the objects before theorizing about the oracle it lacks. The rung $r = 17$ is running as this note is finalized.

8 Postscript II (September 4, 2026): pencils, the frontier, and where the game actually loses

The $r = 16$ rung raised the question whether a stronger *leaf* could lift the substitution game past its exact value 18 (Section 4). We answered it in three steps: an exact pencil oracle, a dump of the adversary’s winning frontier, and exact-rank verdicts on that frontier.

8.1 An exact pencil-rank oracle over $\mathbb{F}_2$, and three textbook facts that fail

A state whose thinnest side has dimension 2 is a matrix pencil $sA + tB$, whose rank is determined by its Kronecker canonical form. We implemented the form over $\mathbb{F}_2$ (minimal indices from the nullities of the block-Toeplitz matrices; elementary divisors from the Smith normal form of $A + xB$ and of $xA + B$ over $\mathbb{F}_2[x]$) and computed the rank of each indecomposable block by brute force, $\min_X \text{rank } X + \text{rank}(I+X) + \text{rank}(C+X)$, for L_ε with $\varepsilon \leq 4$ and regular blocks of degree ≤ 5 . Three statements that hold over algebraically closed fields fail over $\mathbb{F}_2$ (each a brute-force fact, checked against 312 random small pencils):

1. $\text{rank } L_3 = 5$, not Ja’Ja’s $\varepsilon + 1 = 4$: L_ε is the multiplication of a linear form by a polynomial of degree $\varepsilon - 1$, which needs $\varepsilon + 1$ evaluation points, and $\mathbb{F}_2 \cup \{\infty\}$ has three.
2. Regular blocks exceed $D + [\text{non-split}]$: an irreducible cubic contributes $D + 2 = 5$, $(x + 1)^4$ and an irreducible quartic contribute 6.
3. *Block additivity fails*: $(x + 1)^2 \oplus N_2$ has rank 5, not $3 + 3$. The block sum is therefore unsound as a lower bound over $\mathbb{F}_2$; it overshoot the true rank in 5 of 312 random pencils.

The leaf actually installed is the sound cross-block substitution bound, $\sum_i s_i + \max_i e_i$ (s_i the block’s algebraically-closed value, e_i its $\mathbb{F}_2$ excess), exact by brute force for $mn \leq 16$; it is tight on 308 of the 312 pencils.

Observation 2 (pencil leaves lift nothing). *With the sound pencil leaf the Koszul+ $\{A, B\}$ game still has exact value 18 ($k = 19$ refuted exhaustively, 12,965,832 nodes, 6.2 h). With the optimistic block-additive leaf — unsound, hence an upper estimate of any pencil leaf’s power — $k = 19$ is again refuted (11,074,873 nodes, 7.4 h).*

8.2 The adversary’s winning frontier

We instrumented the game to dump every refuted state (all prover moves failed) with its depth, quotient dimensions, target k , leaf value, and — when some side has dimension ≤ 4 — the quotient tensor itself. A 3-hour sample of the $k = 19$ refutation (6.9M of its ~ 13 M nodes) produced 822,622 refuted states, 783,417 distinct tensors, at depths 8–13; side C is never killed in the $\{A, B\}$ game, and the target exceeds the leaf by 1 (408K states) or 2 (337K), rarely more.

For each dimension class we sampled 24 states, compressed each to its concise form, and asked a SAT solver whether $\text{rank} \leq k - 1$: satisfiable means the adversary’s kills genuinely pushed the rank below the target (*genuine*); unsatisfiable means the true rank meets the target

and the leaf is short (*leaf-limited*). Plain cadical at 60s decided 218 of 331: 213 genuine, 5 leaf-limited. A four-arm A/B on the 113 timeouts at 300s per call (cadical; cadical with a hand lex-leader ordering of the products on their (a, b) bits; kissat; and the in-tree `hydra_satsuma` pipeline, Cook/XOR-GE/satsuma/kissat) decided 47 more, all leaf-limited, 46 of them by the lex-leader arm alone; the other three arms decided 3–4 each, and no two arms disagreed. The measured lever for exact rank by SAT on these instances is the product-permutation symmetry, broken by hand; neither kissat nor satsuma’s automatic breaking converted it into speed within the budget.

class (dims, target)	population	genuine	leaf-lim.	open	est. leaf-lim.
$(2, \cdot, 9), k = 6-10$	96,878	119	1	0	~1,500
$(3, 3, 9), k = 7$	34,749	16	8	0	~11,600
$(3, 4, 9), k = 8$	163,835	12	12	0	~81,900
$(3, 5, 9), k = 9$	136,424	16	8	0	~45,500
$(4, 4, 9), k = 9$	137,424	10	14	0	~80,200
$(3, 6, 9), k = 10$	20,766	13	9	2	~7,800
$(4, 5, 9), k = 10$	174,556	7	0	17	unknown
$(3, 7, 9), (4, 6, 9), (4, 7, 9), k = 11-12$	18,785	21	0	46	unknown
every side ≥ 5	37,904	—	—	—	no tensor recorded

Columns three to five are sample counts (24 per class); the last column extrapolates proportionally and carries a sampling error of roughly $\pm 20\%$ of the class.

8.3 Reading

The game loses in two different ways at two different scales. At the small end — every 2-dimensional side, and the smallest 3-dimensional states — the refutations are genuine: the adversary’s kills lower the rank below the target, and no leaf can change that; this is why both pencil leaves lifted nothing. At the mid-size frontier — concise tensors of roughly $3-4 \times 6-7 \times 6-7$ at targets 8–10 — the decided states are leaf-limited 47 times out of 48: their true rank meets the target and Koszul is one short. Over the sampled classes the extrapolation is about 414K genuine, 228K leaf-limited and 141K open states, the open mass sitting in $(4, 5, 9)$ at $k = 10$ and above, where lex-leader cadical at 300s decided nothing, and in the 37,904 states with every side of dimension ≥ 5 , which the dump did not record.

Two consequences. A leaf that certifies 8–10 on those mid-size shapes would invalidate the adversary’s refutations there; whether the root then reaches 19 is a new question, because the adversary may have other replies. And no such leaf is a table or a SAT call: the sample holds 783K distinct tensors, lex-leader SAT needs 0.3–267s each, and the number of GL -orbits of concise $3 \times 6 \times 7$ tensors over $\mathbb{F}_2$ is about 2^{26} . What is needed is a cheap bound that closes a gap of exactly one on these specific shapes. The natural test set is the class $(3, 4, 9)$ at $k = 8$: 164K states, half leaf-limited on the sample, compressed shape about $3 \times 4 \times 5$; its 47 refuted instances are archived with their CNFs in the repository.

The probe as a leaf, and the lattice test. We then tried the one leaf we had that is not in the flattening class: the pooled deep probe itself (concise residuals, kills on all three sides, trimmed Koszul, memo, bounded depth). On the 331 sampled frontier states it is sound (0 certifications on the 214 genuine states) and certifies 20 of the 52 leaf-limited and 33 of the 65 open states at depth 6, most in under a second; depth 10 adds 9 more and exhausts the rest.

Installed in the game as a budgeted, memoized leaf at states with a side of dimension ≤ 4 , it was called 2.77 million times and certified 57% of those states — and $k = 19$ was refuted anyway (8.28M nodes, 8.9 h): the adversary routes the refutation through the genuine half of the frontier and the 3-dimensional-side classes. The exact value 18 now stands under four rule sets. Finally, on Wang’s own lattice — all 14 codim-2 and 68 codim-3 orbits, quotients $7 \times 9 \times 9$ and $6 \times 9 \times 9$ — the probe at depth 8 certifies 0 of 164 targets, failing even his table values: his DP proves 17 and 18 by full-depth substitution memoized over the orbit lattice, while a bounded probe from target 18 must prove 10–11 on residuals with two 9-dimensional sides, which no leaf in this class can do. The route to 21 through a stronger leaf is closed; what remains is a change of move set or an argument outside substitution.

9 Postscript III (September 6, 2026): the code bound, the $r = 17$ verdicts, and a standing assessment

The code bound. In a decomposition $\mathcal{T} = \sum_{i \leq r} a_i \otimes b_i \otimes c_i$, the slice $\mathcal{T}(\varphi, \cdot, \cdot)$ of a covector φ on side A is the sum of the rank-one matrices of the products with $\varphi(a_i) = 1$, so $|\{i : \varphi(a_i) = 1\}| \geq \text{rank } \mathcal{T}(\varphi, \cdot, \cdot)$. Over $\mathbb{F}_2$ the map $\varphi \mapsto \{i : \varphi(a_i) = 1\}$ is linear, so the products’ coefficient vectors form a k -dimensional binary linear code in $\mathbb{F}_2^k$ whose weights are bounded below by slice ranks. With m_c the number of products of coefficient type $c \in \mathbb{F}_2^k \setminus 0$,

$$\text{rank}(\mathcal{T}) \geq \min \left\{ \sum_c m_c : \sum_{c: \langle c, \varphi \rangle = 1} m_c \geq \text{rank } \mathcal{T}(\varphi, \cdot, \cdot) \text{ for all } \varphi \neq 0 \right\},$$

an integer program with $2^k - 1$ variables, maximized over the three sides; its LP relaxation is made rigorous by rationalizing the dual and computing the value-to-load ratio in exact integers. The coset bound of Section 4 is the special case “every slice has rank ≥ 3 ”. At the root the bound is 16 — above Koszul’s 14 and the coset bound’s 13 — with the exact certificate 4032/257; it is 16 on all 211 level-1 residual orbits (the $r = 16$ rung, proved by search in 65 minutes, follows from 263 LPs), 16 on 94% of a root’s 256 folds and 15 on the rank-one killers, 13–15 on Wang’s codim-2/3 orbits against his 17–19, and it certifies 5 of 117 frontier states. Known values are never exceeded ($\langle 2, 2, 2 \rangle$: 6; $\langle 2, 2, 3 \rangle$: 9; $\langle 2, 3, 3 \rangle$: 11). Its strength is at the top of the search tree, the opposite profile from the deep probe.

$r = 17$ verdicts. Installed as a leaf in the pooled probe (HiGHS, 0.1–0.4 s per LP), the code bound collapses the top of the tree: at $r = 17$ on root 0 it certifies 24 of 28 level-1 representatives, 491 of 512 level-2 and 1,148 of 1,240 level-3 nodes, and root 0 — unresolved after 70 core-hours on August 31 and after two further 48-core-hour attempts — is proved in 2.75 h (33 core-hours; 74M tasks). A second root takes 26 minutes. The rung $r = 17$ (which proves rank ≥ 18) is therefore 2,500–7,000 core-hours: 9–24 days on this machine, an afternoon on a thousand cores. Wang’s DP proves ≥ 20 here in under an hour.

Standing assessment. On the rank bound itself, nothing in this program competes with Wang’s DP: it reaches 20 in under an hour where our independent ladder reaches 17 now and 18 in days, and the reason is structural (his memo shares work across a group of order 168^3 ; ours across stabilizers of order 2–768). We are not on track to 21. The substitution game’s value 18 has held under four rule sets, every leaf that beat flattening was routed around by the

adversary, and the deep probe cannot reproduce his codim-2/3 values, let alone lift them. What the program has produced are results about the problem’s structure — the two fold lemmas, the corrected orbit set, the $\mathbb{F}_2$ pencil oracle with its three corrections to textbook values, the frontier map, and the code bound — each backed by an exhaustive run or an exact certificate, and none a step toward 21. The one candidate still untested is an argument of a different shape: a single integer program over the products’ coefficient types whose constraints are *all* of Wang’s verified lattice values at once ($\sum_{c \notin U} m_c \geq t_U$ for every killed subspace U), which his DP uses only one child at a time; it reproduces his 20 from the codim-1 values by integrality alone, and whether the codim-2 and codim-3 values force 21 is a computation, reported below if it is run.

10 Convergence, the frontier at $r = 20$, and outlook

Three structurally different programs — mixed-side substitution inside Wang’s DP, the exhaustive game, and certified scheme search — terminate at the same gap: a cheap ($\lesssim$ ms) rank lower-bound oracle of strength 15–19 for *arbitrary* tensors (lattice-restricted tensors have Wang’s table; residuals do not). The barrier section explains why the gap is stubborn: that strength range spans the rank-method ceiling, so the oracle must itself be amplification-shaped.

For the program’s actual goal, the frontier is sharper than any oracle question. Subadditivity plus Wang’s certificate gives $\text{rank}(T \oplus m_1 \oplus \cdots \oplus m_k) \geq 20 - k$ for free, which trivializes every rung $r \leq 19$; proving $R_{\mathbb{F}_2} > 20$ requires certifying $21 - k$ on *scheme-consistent* prefixes — beating subadditivity by exactly one, which at $k = 0$ is the theorem itself. The +1 must come from knowledge about minimal decompositions of $\mathcal{T}$ specifically. Two routes: (a) prove and iterate the Rank Law (killer varieties as case-split rules — the game’s critical directions may thin as prefixes accumulate); (b) decomposition-structure theorems in the tradition of de Groote’s uniqueness and Burichenko’s symmetries, mined from the known scheme corpora and imported as WLOG prunes. On the constructive side, the corridor experiment established that prefix descent has no usable guidance signal (flattening is capped and constant; true schemes’ intermediates are low-rank, not low-popcount), so the rank-22 search correctly factors as SLS near-misses (native-ANF Brent engine) finished or refuted by the exact endgame engine, which refutes junk prefixes in milliseconds.

All measurements, certificates, logs and the killer dataset are in the repository; every number above is reproducible from committed artifacts.

—

Acknowledgment of method. The program’s productivity traces to two disciplines applied without exception: gate every new rule on known exact values (five unsound artifacts caught), and distrust suspicious nulls until the treatment is proven applied (three vacuous experiments unmasked). The instruments built for soundness — the activity-deficit counters, the killer dump — produced the mathematical discoveries.